

HERMOD processing suite

Authors: Joakim Möller <joakim.moller@molflow.com>
Donal Murtagh <donal.murtagh@chalmers.se>
Joachim Urban <joaurb@chalmers.se>

Version: 0.2

Date: 2010-10-01

Abstract

Hermod is a part of the Odin processing chain - automating level1b data to level2 data. This document describes the installation, configuration and function of the program suite. In short Hermod prepares requirements such as external data files and launches a Qsmr session when all requirements are met.

Contents

1	Level2 processing chain - Hermod	3
1.1	Overview	3
2	Required dependencies - Installation and configuration	5
2.1	Installation of the Database	6
2.2	Torque configuration	6
2.2.1	Torque client configuration	6
2.2.2	Torque server configuration	6
2.2.3	Torque starting and stopping	7
2.3	Maui configuration	7
3	HERMOD	9
3.1	Overview	9
3.2	Package details	9
3.3	HERMOD Installation	9
3.3.1	Developers installation	9
3.3.2	Tests	10
3.3.3	Developing and creating eggs for production	10
3.3.4	Installation in Production environment	11
3.4	Data model	11
3.4.1	level1	11
3.4.2	status	11
3.4.3	Downloading level1 files	12
3.4.4	Finding scans available for processing	12
3.4.5	Queuing and execution	13
3.4.6	Processing	13
4	Troubleshooting	15
5	Appendix A - MySQL Create script	17
6	Appendix B - MySQL Table layout	19
7	Appendix C - Torque server settings	21
8	Appendix D - Maui configuration	23

1 Level2 processing chain - Hermod

The processing chain program suite is a set of python modules that provides an information system that makes it possible to track every single Odin Level1 file and choose a suitable processor to make higher level data i.e. Level2 data.

Hermod is a part of the Odin processing chain making high level information from calibrated satellite data to very high level data ie. Human understandable data and possibly data collected and aggregated over longer time periods.

1.1 Overview



The Hermod suite are written mostly in [Python](#)¹ and small part of the code is written in C with Python's C-api to extend Python's capabilities to interact with different tools in the Processing chain.

Meta data from calculations made by Qsmr is stored in a database and data is stored in files at the file system.

Hermod is the collection name for the processing chain. The name Hermod was chosen after one of the sons of Odin - Hermod known for his speed.

The Hermod processing system can be seen as a set of scripts that glues Qsmr's calculations and its results into the database. Hermod also uses those result to find out what data is missing or what can be calculated for the moment ie. all prerequisites for starting Qsmr calculations are resolved. Hermod also serves other automated systems like IASCO model with data.

2 Required dependencies - Installation and configuration

The Odin processing chain and Hermod make use of third party software. They are all based on some type of open source license like GNU GPL or BSD license.

Hermod is built to run on Ubuntu Linux 10.04 (server version) but may work on different Ubuntu versions as well as other POSIX OS:es probably even on windows.

Hermod needs other components to work properly:

Python ¹:

Hermod's core is implemented in Python 2.6. But other versions may also work. Hermod uses a lot of external packages see below.

MySQL ²:

Relational database to manage meta data. Database installation for this project is discussed in [Installation of the Database](#)

Torque ³:

Torque is a Cluster Resource Manager. Documentation and detailed installation instructions can be found at [Torque](#) ³ documentation pages. Site specific configuration will be discussed in [Torque configuration](#) section.

Maui ⁴:

The Cluster Scheduler only site specific setup will be noted in [Maui configuration](#)

On a Ubuntu 10.04 LTS machine for developing Hermod the following apt-packages need to be installed:

```
gfortran
libatlas-base-dev
libblas-dev
libfreetype6-dev
libfuse-dev
libgeos-dev
libhdf4g-dev
libjasper-dev
libjpeg62-dev
libmysqlclient-dev
libpng12-dev
libtorque-dev
openssh-server
pkg-config
python-dev
python-matplotlib
python-numpy
python-scipy
python-virtualenv
python-virtualenv
subversion
torque-dev
ubuntu-dev-tools
vim-nox
zlib1g-dev
```

2.1 Installation of the Database

Configuration of the database is minimal - Hermod works fine on a standard apt installation of the package. But further tuning may increase performance significantly. See [Appendix A - MySQL create script](#) and [Appendix B - MySQL Table layout](#) for database and table layout.

On the nodes at least `libmysqlclient` needs to be installed.

2.2 Torque configuration

Two types of Torque installations are required - one server installation and several client installations on each node in the cluster. The server installation manages the queueingsystem and needs to know about all clients (compute nodes) in the cluster. The clients does only need to know about the server.

2.2.1 Torque client configuration

A site-specific installation script `/misc/apps/torque-package-mom-linux-x86_64.sh` provided all configuration needed at the client. But some additional configuration is needed to provide the per session temp directory.

The following script makes all steps in the installation process.

```
#!/bin/bash
# A script to install, prepare and start a node
# run as root

aptitude purge torque-mom torque-client -y
sh /misc/apps/torque-package-mom-linux-x86_64.sh --install
cp /misc/apps/prologue.user /var/spool/torque/mom_priv/
cp /misc/apps/epilogue.user /var/spool/torque/mom_priv/
ldconfig
pbs_mom
```

An important part of the processing system is the scripts at the client that creates a temporary directories before a processing starts and removes it when processing is finished. These scripts runs whether or not the processing was successful or not.

2.2.2 Torque server configuration

A site-specific installation script `torque-package-server-linux-x86_64.sh` installs binaries and libraries and some basic configuration. Editing configuration files to reflect connected nodes and their capabilities is necessary.

The file `/var/spool/torque/server_priv/nodes` defines the compute nodes:

```
glass np=8 hermod node x86_64
sard np=2 hermod node x86_64
...
```

The attributes `hermod`, `node` and `x86_64` specifies different capabilities en each node. 'x86_64' tells us the architecture on the node is 64 bits. 'hermod' states that hermod, Qsmr and Q-pack is installed and works correctly. The last attribute shows us the computer is a node with no other users than the torque queue operates the computer. 'desktop' would state it is a workstation with human users.

Some additional settings can be done through torque's configuration program `qmgr`. A printout of Torque server settings generated with `qmgr -C 'print server'` can be found in [Appendix C - Torque server settings](#).

2.2.3 Torque starting and stopping

There are currently no system V init scripts implemented. Starting and stopping server and nodes is manual. There is no problem shutting off a node before the server but the running job at the node will be killed. If server is stopped the current queue will be saved and the current running jobs at the moms will continue. When server is started again moms will report their finished jobs.

start server at morion:

```
$ /usr/local/sbin/pbs_server
```

start moms at nodes:

```
$ /usr/local/sbin/pbs_mom
```

stop moms at nodes:

```
$ /usr/local/sbin/momctl -s
```

stop server at morion:

```
$ /usr/local/bin/qterm -t immediate
```

2.3 Maui configuration

The main configuration file can be found at `morion.rss.chalmers.se`.

`/usr/local/maui/maui.cfg`

Full configuration file can be found in [Appendix D - Maui configuration](#). This setup restrict one user to take all resources at once enforcing Odin processing always have at least a minimum of processor available but also giving users access to the queue.

start the scheduler:

```
$ /usr/local/maui/sbin/maui
```

stop the scheduler:

```
$ /usr/local/maui/bin/schedctl -k
```


3 HERMOD

3.1 Overview

Hermod is a program suite written in Python that wraps around QSMR and inserts meta data in to the SMR database. Hermod runs regularly and decides when to run QSMR according to information Hermod can find in the SMR Database. Hermod provides a fully automatic processing system for processing data from Level1 data to Level2 data.

3.2 Package details

Hermod is divided into several smaller entities that provide specific functionality. The current status of the source code is still in a form of transition from one package to more and smaller sub packages.

odin.hermod

The odin.hermod package is the package which is responsible for the information and bookkeeping parts of hermod i.e keep track of file transactions, file dependencies and finally submitting jobs to the queuing system

odin.config

The odin.config is more or less a configuration package Hermod and lasco shares this package

odin.iasco

This runs and manages the IASCO model.

3.3 HERMOD Installation

For the moment hermod is running from the development source i.e. from the directory `~odinop/hermod_jm` for Ubuntu 10.04 and `~odinop/hermod_glass` for 9.08 this directory is checked out from svn. This is not by any means the ideal way to maintain a piece of software. This is a temporary solution.

Best way to continue development is to separate development and production. First all processing nodes and servers in the system need to have the same OS version (Ubuntu 10.04 LTS). Using the same OS makes it possible to run Hermod from on single installation shared by NFS.

A set of compiled hermod packages exists in `/misc/apps/odinsite`. The installation of hermod is controlled by `zc.buildout` installation. This way buildout pins down the specific version of each dependency package.

3.3.1 Developers installation

The source of hermod is available at [Chalmers' Subversion repository](#) ⁵. A developers installation is a isolated installation which is running in its own environment - from here it's possible to run unit test and other functional testing. With a correct `.hermod.config` and `.hermod.config.secret` it's possible to connect to the database or PDC. (Future work: It would even be possible to have a standalone database with a small data set to run off site tests).

Once developers build packages in the development environment described above. Packages selected to be "released" are copied to `/misc/apps/odinsite`.

To have a copy of hermod running locally for development do the following:

```
svn co http://svn.rss.chalmers.se/svn/odinsmr/hermod/trunk hermod
virtualenv -p /usr/bin/python2.6 hermod
cd hermod
bin/python2.6 bootstrap.py
bin/buildout
```

3.3.2 Tests

To ensure quality and to simplify for other developers. Unittests are written for some parts of Hermod. Unittests makes sure that the tested function or procedure does what it's meant for.

Example: How to run a full unittest suite for a package.

```
.. src/odin.ecmwf$> ../bin/odinpy setup.py test
```

Example: how to run a specific testcase.

```
.. src/odin.ecmwf/odin/ecmwf/tests$> ../../../../../../bin/odinpy zpt2_create_test.py
```

3.3.3 Developing and creating eggs for production

Once you have a [Developers installation](#) you can change or correct Hermod's behavior. If you want to deploy your changes you have to change the version variable in the setup.py file. When tests are ok - commit your changes to svn. If you don't have any test for your code - consider to add a test to cover your code.

To build installable eggs use the python interpreter created with buildout.

```
$ cd <devel>
$ cd src/odin.hermod
$ ../../bin/odinpy setup.py bdist_egg
```

To install an egg in a production environment.

```
$ cd <production>
$ bin/easy_install -U -f <where the eggs are> odin.hermod
```

1. make changes
2. run tests
3. commit
4. build eggs
5. install in production environment.

This is the cycle I use when developing for Hermod.

1. check out from svn
2. create environment
3. develop
4. run tests
5. check in
6. release
7. deploy
8. remove development files

Step 1-2 is normally done once.

Step 3-4 iterates many times. Step 3-5 when tests are successful.

Step 8 is seldom made.

3.3.4 Installation in Production environment

Released packages can be installed in the production environment by using Python's `easy_install` utility.

Packages can be installed either into the system environment or into a virtual python environment. (Preferably the virtual environment to not clutter the system installation)

```
virtualenv <dir>
cd <dir>
bin/easy_install -f /misc/apps/odinsite \
    odin.config \
    odin.hermod \
    odin.iasco
```

Later on updates can be installed by:

```
cd <dir>
bin/easy_install -f /mist/apps/odinsite -U \
    odin.hermod
```

3.4 Data model

The database consists of a number of loosely connected tables with records (rows) describing meta data about satellite measurement or metadata of files stored on disk.

The Hermod data model is pretty simple. All tables are 'knitted' together with a 'id' field. For example in the 'level1'-table the logical key that identifies each row is the fields 'orbit','calversion' and 'freqmode'.

3.4.1 level1

This table contains the metadata from the process of producing 'Level 1' data at Onsala. One orbit of Odin corresponds to at least 2 rows in the database, one for each combination of freqmode, calversion and backend.

```
id -> orbit, calversion, backend, freqmode -> 'records in level1'
```

3.4.2 status

Some error messages from the level0 to level1 process are captured in this table.

```
id -> status,errmsg
```

The 'id'-field is included in the 'level2'-table to make it possible to find all level2 products derived from a 'level1' record.

level2:

```
id, fqid, scanno -> 'records in level2-table'
```

level2files:

```
id, fqid -> 'records in level2files-table'
```

3.4.3 Downloading level1 files

Hermod searches the database to find new files available on PDC but not in the local file storage.

```
select l1.id,l1.filename,l1.logname
from level1 l1
join status s on (l1.id=s.id)
left join level1b_gem l1bg on (l1.id=l1bg.id)
where s.status and (l1bg.id is null or l1bg.date<l1.uploaded)
and s.errmsg=' ' and l1.calversion in (6,7);
```

3.4.4 Finding scans available for processing

To find new orbits in the database that has not already been processed to a level2 file.

```
select distinct l1.id,l1.back backend,l1.orbit orbit,v.id fqid,
v.qsmr version, l1.calversion,a.name,v.process_time
from (
select orbit,id,substr(backend,1,3) back,freqmode mode,
calversion from level1
join status using (id)
join level1b_gem l1g using (id)
where status and l1g.filename regexp ".*HDF"
and not locate(', ',freqmode)
union (
select orbit,id,substr(backend,1,3) back,
substr(freqmode,1, locate(', ',freqmode)-1) mode,
calversion from level1
join status using (id)
join level1b_gem l1g using (id)
where status and l1g.filename regexp ".*HDF"
and locate(', ',freqmode)
)
union (
select orbit,id,substr(backend,1,3) back,
substr(freqmode from locate(', ',freqmode)+1) mode,
calversion from level1
join status using (id)
join level1b_gem l1g using (id)
where status and l1g.filename regexp ".*HDF"
and locate(', ',freqmode)
)) as l1
join versions v on (l1.mode=v.fm)
join Aero a on (v.id=a.id)
left join level2files l2f on
(l1.id=l2f.id and v.id=l2f.fqid and v.qsmr=l2f.version)
left join statusl2 s2 on
(l1.id=s2.id and v.id=s2.fqid and v.qsmr=s2.version)
where v.active and l2f.id is null and l1.calversion=6
and (proccount is null or proccount<4)
order by orbit desc,fqid
```

3.4.5 Queuing and execution

A "job" is defined from the look up in the previous section. Information about the processing is sent to a queue for later execution. The Resource system that handles the queue and the execution nodes in the computing cluster (`glass`, `larimar`, `titanite`, ...) is [Torque](#)³.

Basically the "job" is a shell script sent to another machine for execution.

The script `run_processor` puts the shell script in queue with different input parameters to run on the compute nodes.

3.4.6 Processing

The `hermodprocessor`-script executes the main-function in `odin.hermod.processor-module`. This module looks in the database to find level1b records which not have as many corresponding level2 records as hermod expects.

When Hermod detects a job to run - Hermod sends a wrapped Qsmr job to the processing cluster and collects the results and puts them in the database and the file system.

4 Troubleshooting

A job is stale - showing negative time with qstat:

The execution service on the node is probably dead. Use *qstat -rn* to see what node the job runs on, also note the jobnumber. Log in as root at the stale node. Start the mom by *pbs_mom*. When the mom is started launch *momctl -c <jobnumber>* to clear the nodes status.

5 Appendix A - MySQL Create script

6 **Appendix B - MySQL Table layout**

7 Appendix C - Torque server settings

```
#
# Create queues and set their attributes.
#
#
# Create and define queue batch
#
create queue batch
set queue batch queue_type = Execution
set queue batch resources_default.nodes = 1
set queue batch resources_default.walltime = 01:00:00
set queue batch enabled = True
set queue batch started = True
#
# Create and define queue new
#
create queue new
set queue new queue_type = Execution
set queue new resources_default.nodes = 1
set queue new resources_default.walltime = 01:00:00
set queue new enabled = True
set queue new started = True
#
# Create and define queue new
#
create queue rerun
set queue rerun queue_type = Execution
set queue rerun resources_default.nodes = 1
set queue rerun resources_default.walltime = 01:00:00
set queue rerun enabled = True
set queue rerun started = True
#
# Set server attributes.
#
set server scheduling = True
set server acl_hosts = morion
set server managers = root@morion.rss.chalmers.se
set server operators = root@morion.rss.chalmers.se
set server default_queue = batch
set server log_events = 511
set server mail_from = adm
set server query_other_jobs = True
set server scheduler_iteration = 600
set server node_check_rate = 150
set server tcp_timeout = 6
set server mom_job_sync = True
set server keep_completed = 300
set server auto_node_np = True
set server next_job_number = 18315
```


8 Appendix D - Maui configuration

The only configuration file is in `/usr/local/maui/maui.cfg`:

```
# maui.cfg 3.3

SERVERHOST          morion
# primary admin must be first in list
ADMIN1              root e0joakim jo
ADMIN2              donal odinop
ADMIN3              all

# Resource Manager Definition

RMCFG[base] TYPE=PBS

# Allocation Manager Definition

AMCFG[bank] TYPE=NONE

# full parameter docs at http://supercluster.org/mauidocs/a.fparameters.html
# use the 'schedctl -l' command to display current configuration

RMPOLLINTERVAL      00:00:30

SERVERPORT          42559
SERVERMODE          NORMAL

# Admin: http://supercluster.org/mauidocs/a.esecurity.html

LOGFILE             maui.log
LOGFILEMAXSIZE      10000000
LOGLEVEL            3

# Job Priority: http://supercluster.org/mauidocs/5.1.jobprioritization.html

QUEUETIMEWEIGHT     1

# FairShare: http://supercluster.org/mauidocs/6.3.fairshare.html

FSPOLICY            PSDEDICATED
FSDEPTH             7
FSINTERVAL          6:00:00
FSDECAY             0.80

FSWEIGHT 10
CREDWEIGHT 100
USERWEIGHT 0
GROUPWEIGHT 0
CLASSWEIGHT 100
SERVICEWEIGHT 1
QUEUETIMEWEIGHT 1
FSCLASSWEIGHT 100
FSUSERWEIGHT 0
```

```

# Throttling Policies: http://supercluster.org/mauidocs/6.2throttlingpolicies.html

# NONE SPECIFIED

# Backfill: http://supercluster.org/mauidocs/8.2backfill.html

BACKFILLPOLICY          FIRSTFIT
RESERVATIONPOLICY       CURRENTHIGHEST

# Node Allocation: http://supercluster.org/mauidocs/5.2nodeallocation.html

NODEALLOCATIONPOLICY     MINRESOURCE

# QOS: http://supercluster.org/mauidocs/7.3qos.html

# QOSCFG[hi]  PRIORITY=100 XFTARGET=100 FLAGS=PREEMPTOR:IGNMAXJOB
# QOSCFG[low] PRIORITY=-1000 FLAGS=PREEMPTTEE

# Standing Reservations: http://supercluster.org/mauidocs/7.1.3standingreservations.html

# SRSTARTTIME[test] 8:00:00
# SRENDTIME[test]   17:00:00
# SRDAYS[test]       MON TUE WED THU FRI
# SRTASKCOUNT[test] 20
# SRMAXTIME[test]    0:30:00

# Creds: http://supercluster.org/mauidocs/6.1fairnessoverview.html

USERCFG[DEFAULT]        FSTARGET=20 MAXJOB=10
USERCFG[odinop]         FSTARGET=50 MAXJOB=50
# USERCFG[john]          PRIORITY=100 FSTARGET=10.0-
# GROUPCFG[staff]        PRIORITY=1000 QLIST=hi:low QDEF=hi
#CLASSCFG[batch]         FLAGS=PREEMPTTEE
CLASSCFG[batch]          FLAGS=PREEMPTTEE PRIORITY=10000
# CLASSCFG[interactive]  FLAGS=PREEMPTOR
CLASSCFG[batch] FSTARGET=40.0
CLASSCFG[rerun] FSTARGET=20.0
CLASSCFG[new] FSTARGET=40.0

```

-
- 1(1, 2) <http://python.org/>
 - 2 <http://dev.mysql.com/doc/refman/5.1/en/>
 - 3(1, 2, 3) <http://www.clusterresources.com/products/torque/docs>
 - 4 <http://www.clusterresources.com/products/maui/docs>
 - 5 <http://svn.rss.chalmers.se/svn/odinsmr/hermod>